



Bachelor Thesis

Analysis of authoritative DNS infrastructure failures

Ege Girit

December 27, 2022

Reviewer:

Prof. Dr. Christian Rossow
BSc. Jonas Bushart

Contents

1	Introduction	1
1.1	Structure Of The Work	1
1.2	Domain Name System	1
1.3	Motivation	3
2	Structure of the experiment	4
2.1	Parameters To Be Tested	4
2.2	Test Sources	4
2.3	Details Of The Test Environment	5
3	Experiments and Results	7
3.1	Open Resolver Packetloss Experiment	7
3.2	Open Resolver Packetloss Experiment Results	8
3.3	Ripe Atlas Packetloss Experiment	9
3.4	Ripe Atlas Packetloss Experiment Results	10
3.5	Wild Open Resolver Experiment	10
3.6	Wild Open Resolver Experiment Results	10
3.7	Stale Record Support Experiment	10
3.8	Stale Record Support Experiment Results	12
3.9	Stale Record TTL Experiment	13
3.10	Stale Record TTL Experiment Results	13
3.11	Stale Record Send Frequency Experiment	14
3.12	Stale Record Send Frequency Experiment Results	14
3.13	Ripe Atlas Stale Record Experiment	14
3.14	Ripe Atlas Stale Record Experiment Results	14
3.15	DNS Truncation Experiment	14
3.16	DNS Truncation Experiment Results	16
4	Conclusions	17
5	Discussion	18
6	Ethical Considerations	19
	Bibliography	20

1 Introduction

1.1 Structure Of The Work

This bachelor thesis is divided into four main parts. In the first section of this paper, we provide a summary of the Domain Name System (DNS), which forms the foundation of our focus. This overview covers the basics of DNS and sets the stage for the subsequent discussions. Building on this foundation, we explain why we believe it is important to study resolver behavior during an authoritative server failure. In the second part of the paper, we describe the specific characteristics we want to examine, the sources we will use for testing, and the metrics we will use to determine the success of an attack. We also outline the test environment and analysis methods. In the third part, we describe the experiments and present the results of our experiments using visual graphs. We then evaluate these results and draw conclusions about what the tests reveal. Finally, in the fourth part, we discuss the extent to which our results may differ from reality and whether further research is needed in this area.

1.2 Domain Name System

The need for the Domain Name System arose from the need to address the IP addresses of servers, which are difficult for humans to remember, with easy-to-remember names. When a user enters a server's name into their computer or device, DNS is used to resolve the name to the corresponding IP address, enabling the user to connect to the server. Before the Domain Name System, there was only a centralized text file (`hosts.txt`) consisting of static mappings that allowed users to resolve stored web pages to their associated IP addresses. This text file had to be updated after each change and distributed to all computers to provide a mapping between these names and their corresponding network addresses. With the growing number of hosts and websites, it became impractical to manage such a centralized system. The administrative overhead associated with managing every possible domain name on the Internet would be too great, and this central database would not scale well. To address these challenges, DNS designers abandoned the flat naming approach, where the names do not have any internal structure or clear connections to each other, and instead implemented a decentralized model with a hierarchical naming architecture that became the Internet standard in 1986. The official documentation of these standards is described in RFC 1034 and 1035. However, the `hosts.txt` file still exists for various purposes like performance improvement and URL filtering.

1 Introduction

Today we have a globally distributed collection of databases that form a tree structure. System administration is completely decentralized, and responsibility for delegation is handed over to organizations. There are three main functions of DNS:

1. Name resolution
2. Name space
3. Name registration

A namespace refers to the way that domain names are structured and used in terms of what makes a name valid. Name registration is there to guarantee whether the name to be registered is unique. In this work, we will focus on the core function of DNS, which is the name resolution. Several actors play a role in resolving names to their corresponding IP addresses. These actors are:

- The client
- Recursive resolver
- Root name server
- TLD Server
- Authoritative name server

The client is interested in the answer of the name resolution. To start the name resolution process, a client sends a DNS query to a recursive DNS resolver. The DNS resolver then takes the role of the client and starts the querying process by trying to find the answer to the query on behalf of client and returning the results to the client. The recursive resolver acts as a middleman between the client and the DNS hierarchy. The recursive resolver contacts the DNS hierarchy to get the DNS response. The first server that the recursive resolver contacts is the root name server. The root name server sends back the appropriate top-level domain (TLD) server based on the queried host's domain extension to the recursive resolver. The TLD server receives the next request from the recursive resolver and responds with the appropriate authoritative nameserver. The authoritative nameserver stores the DNS resource records that map domain names to IP addresses, and it responds to the recursive resolver's final request with the queried hostname's IP address. If the IP address is unavailable, the nameserver will return an error. A DNS resource record is a type of data that is stored in a DNS database and used to provide information about various aspects of a domain name or IP address. There are several different types of DNS resource records, each with a specific purpose. In this paper we focus on A (Address) record, which maps a domain name to an IP address. Because there are many servers involved in the resolution process, a request from a client can consume a lot of network resources before the resolution succeeds or fails. Therefore, there are optimization mechanisms such as caching, which shorten the resolution time and save resources. The purpose of caching is to temporarily store previously requested data so that future requests for that data can be served faster without contacting all servers in the DNS hierarchy. When a resolver receives a DNS query, it first checks its

1 Introduction

cache to see if it has a recent copy of the requested resource record. If the record is found in the cache, the resolver can return the record to the client without needing to send a query to an upstream DNS server. After a cache entry of the resolver expires, it can be refreshed if a new response to the query is received by the resolver. If the resolution process takes too much time, the resolution times out. A DNS retransmission is the process of sending a DNS query again if no response or an error response is received from the DNS server. DNS retransmissions are typically triggered by a timeout, which occurs when the DNS server does not respond within a certain time frame. The client will typically retransmit the query multiple times before giving up and returning an error to the user. For a DNS request, there can be multiple duplicate requests, and therefore multiple responses to these requests. If a resolver has not received an answer to a query, the resolver can repeat this query, resulting in duplicate queries and possible duplicate responses.

1.3 Motivation

DNS is a fundamental protocol that was not designed with security in mind. There are many types of attacks aimed at exploiting vulnerabilities in DNS. If name resolution does not work properly, there is no alternate protocol for name resolution, and the connectivity between users and the server is compromised. This makes DNS a popular target for hackers. Many systems rely on DNS to remain functional. Therefore, errors or attacks on DNS are of particular importance for everything based on it. If an authoritative name server is unreachable, the client usually has to expect longer waiting times or errors. However, there is a lack of knowledge about the specific behaviors of resolvers during cyber attacks and authoritative server failures. This gap in understanding presents an opportunity for further research.

In this thesis, we investigate the behavior of resolvers in the event of a partial or complete failure of authoritative DNS servers. Through our experiments, we aim to understand whether the behavior of resolvers follows certain patterns and how the defensive mechanisms of DNS work against cyber attacks. By analyzing the results of these experiments, we hope to gain a better understanding of the impact of different parameters on communication with the name server and to predict the behavior of the resolver. This information could be useful in developing strategies to defend against cyber attacks and improve the performance and reliability of DNS.

Similar research has been done in this area where the client's point of view and the caching behaviour of the resolvers is mainly examined in the process of DNS resolution.

2 Structure of the experiment

2.1 Parameters To Be Tested

In our experiments, we want to know whether the simulated cyber attack against the DNS defense mechanisms is successful. For this purpose we measure the following metrics:

- DNS resolution latency, which is the time it takes for a DNS resolver to receive a response from a DNS server in response to a DNS query
- Failure rate of responses to examine the percentage of failed DNS resolutions
- The underlying transport protocol of the DNS packet (either UDP or TCP)
- Rate of stale records to examine whether a resolver is able to answer a client query in case of a name server failure
- DNS retransmission amount to examine how the packetloss amplifies the DNS traffic by resending a DNS query

2.2 Test Sources

Given the large number of resolvers with varying configurations, it is not feasible to test all of the resolvers worldwide and draw general conclusions about resolver behavior. Therefore, we need to design our experiments in a way that provides a representative sample of most types of resolvers. This will enable us to draw meaningful insights about resolver behavior while still being feasible to conduct. The types of resolvers we tested in our experiments are as follows:

- Widely used open resolvers
- Resolvers of uniformly distributed worldwide RIPE Atlas probes
- Wild open resolvers found via DNS scans

For our experiments, we first identified widely used DNS operators. The DNS operators we have chosen are the following, which were obtained from the website publicdns.xyz on December 6, 2022: AdGuard, CleanBrowsing, Cloudflare [2], Dyn, Google [5], Neustar, Norton, Level3, OpenDNS [8], Yandex and Quad9 [9]. These operators represent a diverse range of DNS services that are commonly used by a large number of users. In our experiments, we examined a total of 11 operators and 30 open resolver IP addresses. We carefully selected the IP addresses to ensure that we had at least one resolver IP address for each operator configuration. This allowed us to test the different configurations of each operator. Our experiments on stale records showed that different configurations of

2 Structure of the experiment

the same operator can behave differently when an authoritative DNS server fails. This highlights the importance of considering the specific configuration of an operator when studying resolver behavior.

RIPE Atlas is a global network of probes that help us measure the internet connectivity and reachability in real time. The probes are small devices that are distributed around the world and connected to the Internet through various types of connections. Each probe sends measurements to the RIPE Atlas servers and the collected data is used to provide a detailed view of the Internet’s performance, including metrics such as latency, packetloss, and throughput. The data collected by RIPE Atlas is used to understand the performance and reliability of the Internet. With the use of RIPE Atlas, we are able to examine the performance of the DNS resolvers of the selected probe devices.

2.3 Details Of The Test Environment

For our experiments we use an authoritative name server at Saarland University, that uses BIND9 version 9.16 as an open-source name server software to be able to interact with DNS. We only use one name server with one IP address without the use of any anycast structure or secondary servers. The configuration files of the servers that we used in our experiments can be found in the experiment repository.

iptables is a command-line utility in Linux that allows to configure the tables provided by the Linux kernel firewall and the chains and rules it stores. iptables uses a set of tables, which contain chains of rules that match packets and determine what to do with them. It can be used to set up firewall rules that allow or block traffic based on various criteria, such as the source or destination IP address, the protocol, and the port number. To simulate an attack on our authoritative server, iptable rules are used by configuring it to randomly drop incoming UDP and TCP packets on the server at a specified rate. This allows us to study the effects of a DDoS attack on the performance and behavior of resolvers without causing harm to any real systems.

According to [4], this simulation may not be entirely accurate because it does not take into account queueing delays and is only performed at the final router. However, the impact of this limitation is likely to be minimal due to the dominant influence of packet loss on the overall results [3].

Our authoritative name server acts as both a server, providing DNS answers to a resolver, and a client, creating the DNS queries. To distinguish between client and server network traffic, we use separate IP addresses for the client and the server.

To generate and send DNS requests in our experiments, we use the DNSPython library. All of the DNS queries we send have a default timeout value of 10 seconds. The command-line utility tcpdump is used on our authoritative server to capture and display incoming and outgoing network packets transmitted over the network such as DNS queries and responses. To isolate the recorded network traffic for the respective resolver communications, we use wireshark to filter irrelevant packets, that are not generated by

2 Structure of the experiment

our experiment.

After we have performed the experiments, the filtered capture files are evaluated and the data of interest is extracted and visualised by the use of various plots to gain insights into the resolver behaviors. This helps us to understand how the resolvers respond to different scenarios and identify patterns or trends in their behavior.

Given the low adoption rate of IPv6 in DNS, we used IPv4 for our experiments. This enabled us to focus on a widely used protocol and ensure that our results are relevant to a broad audience.

3 Experiments and Results

3.1 Open Resolver Packetloss Experiment

Since we are unable to perform a real cyber attack on our authoritative name server, we simulate a DDoS attack by using iptable rules to introduce random packet loss for incoming packets on the server at a specified packet loss rate.

DDoS stands for Distributed Denial of Service and it is a type of cyber attack in which a large number of computers, which are often compromised, are used to flood a target network or website with network traffic in an attempt to make it unavailable to users. A DDoS attack can create a high network load on the target server by flooding it with a large volume of traffic, which can make it difficult or impossible for the server to process and respond to legitimate requests. As a result, some network packets may be dropped or may not be sent successfully, leading to a partial or complete disruption of service for legitimate users. For our experiment we use various rates of packetloss and these packetloss rates are 10%, 20%, 30%, 40%, 50%, 60%, 70%, 80%, 85%, 90% and 95%.

To ensure that our measurements accurately reflect the full query lookup process without the benefit of caching, we used a unique query name for each query sent to the resolver. Without caching, the resolver is forced to contact the DNS hierarchy to look up the corres

(how the retransmissions help when packetloss active, paper) Prior research has demonstrated that DNS retransmissions during DDoS attacks can cause legitimate traffic to increase by up to 8 times for servers. When the caches of the resolver are full and DNS retransmissions occur, this can significantly reduce resolution failures for less than the cache lifetimes in the case of 90% query loss [4].

(Query name structure...) (Date time of experiment)

For each packet loss rate, we sent 49 queries to each resolver. This allows us to collect a sufficient amount of data to draw meaningful conclusions about the behavior of the resolvers at different packet loss rates. It is generally not appropriate to send a large number of queries to a DNS resolver in a short period of time, as this can put a strain on the resolver and potentially impact its performance for other users. There is no specific threshold for how many queries is considered "excessive", as it can vary depending on the capabilities of the DNS resolver and the overall level of activity on the network. In this experiment, we introduced a one second delay after each query sending. After Sending 49 queries to a resolver for a packetloss rate, we added a cooldown phase of 10 minutes, where we kept the network capture going to be able to see any delayed dns packets. After the cooldown phase, we continued the experiment with the next packetloss rate, until all the packetloss rates are processed. For a DNS query, we calculate the latency

3 Experiments and Results

of the DNS resolution in our experiments by calculating the time difference between the first DNS query and the first DNS response to it in case of duplicate responses.

3.2 Open Resolver Packetloss Experiment Results

To visualise the observed latencies, we use box plot and violin plot, which allows us to see the range as well as the density of the latencies. The thicker part of the plot represents the denser regions of the data, while the thinner part represents the sparser regions of the data.




Figure 3.1: Latency violin plot

If we compare the latencies at a packet loss rate of 0% to a packet loss rate of 95%, we see that the mean latency increases by a factor of 28, while the median latency increases by a factor of 10. Between a packet loss rate of 0% and 50%, the median latency increases by a factor of 2.5, while the mean latency increases by a factor of 5.6. In terms of the worst-case scenario, considering the maximum observed latencies, we see that the latency increases by a factor of 9 between a packet loss rate of 0% and 50%, and by a factor of 32 between a packet loss rate of 0% and 95%. Since the median is more resistant to outliers and the mean is heavily influenced by extreme values, it is generally more reliable to use the median value when performing the latency analysis. In our experiments, we observe many individual outliers which significantly influences the mean values due to their extreme values. Although the graph indicates that latencies

3 Experiments and Results

decrease between packet loss rates of 85% and 95%, there is a significant difference in the amount of latency observed at these rates. Specifically, we see that there are 334 latencies at a packet loss rate of 85%, while there are only 92 latencies at a packet loss rate of 95% due to the increasing failure rates. Additionally, the failure rate is 14% at a packet loss rate of 50%, and it increases to 86% at a packet loss rate of 95%.




Figure 3.2: Failure rates

Upon examining the number of unanswered queries, we noticed that up to a packet loss rate of 60%, there were only a few unanswered queries for each packetloss rate. However, from a packet loss rate of 60%, the number of unanswered queries began to increase significantly. Additionally, we observed that at a packetloss rate of 70% and higher, the number of queries that were not visible on the authoritative name server side increased significantly. When we examine the violin plot for query retransmissions, we observe that a single query can result in up to 19 retransmissions, which is the case for 60% packetloss rate. From a packet loss rate of 60% to 95%, the number of retransmissions for a single query begins to decrease. It is possible that some resolvers may recognize that there is an issue with the authoritative name server and give up on resolving the query earlier, resulting in a response of SERVFAIL.

3.3 Ripe Atlas Packetloss Experiment

We conducted our packet loss experiment with the RIPE Atlas measurement system over two days, due to the daily credit limit. On the first day, we tested packet loss rates of 40%, 60%, 70%, 80%, 90%, and 95%. On the second day, we tested packet loss rates of 0%, 10%, 20%, 30%, 50%, and 85%. In total, we conducted the experiment using 830

3 Experiments and Results

uniformly distributed probes worldwide. The Ripe Atlas probe query timeout value is set to 30 seconds to prevent the report from indicating "no answer" when latency is less than 30 seconds. This ensures that the report accurately reflects the latency of queries, even when they are processed relatively slowly. Similar to the packet loss experiment with open resolvers, we used unique query names for each probe query to prevent caching. This helps to ensure that the results of the experiment accurately reflect the performance of the system being tested, rather than being influenced by cached responses.

3.4 Ripe Atlas Packetloss Experiment Results

Errors or issues with the configuration of the RIPE Atlas device, or policies implemented by the respective resolver, may be responsible for the occurrence of REFUSED answers. These types of issues can cause DNS queries to be rejected or blocked, resulting in the REFUSED response. We discarded probes from the experiment, that sends constant REFUSED or SERVFAIL as RCODES (Only x% of the queries).

3.5 Wild Open Resolver Experiment

We expanded our packet loss experiment to include a larger number of resolver IP addresses. To do this, we scanned the IPv4 address space and identified all IP addresses that appeared to be functioning as resolvers. In total, we identified 6 million IP addresses. We then conducted the same packet loss experiment using these resolvers, but modified the procedure so that we only sent one query to each identified resolver IP due to the large amount of IP addresses. This approach allowed us to test a larger number of IP addresses, but with fewer queries per resolver. It's worth noting that the number and IP address of the identified resolvers may change frequently, so it's possible that some of the resolvers may stop responding during the course of the experiment.

3.6 Wild Open Resolver Experiment Results

3.7 Stale Record Support Experiment

A DNS request from a client can consume a lot of network resources until the resolution succeeds or fails. Therefore, there are optimization mechanisms such as caching, which shorten the resolution time and save resources. The purpose of caching is to temporarily store previously requested data so that future DNS requests for that data can be served faster without contacting all servers in the DNS hierarchy. The resolver stores the cached entries and these cached entries have a TTL (Time to live) value. This value represents the number of seconds that the resolver should keep the entry in its cache. However, in practice, the resolver may alter this TTL value based on its implementation, resulting in the entry being kept for more or less time[49 Unbound Documentation <https://unbound.docs.nlnetlabs.nl/en/latest/>]. According to the studies made by ...,

3 Experiments and Results

the resolvers usually shortens the TTL time rather than keeping the data more than the TTL value. A lower TTL value means that the lookups for the resolutions are needed to be performed more frequently. The resolvers can also manually clear the cache before the TTL value expires for various purposes such as to free up storage space.

If the resolver detects that the authoritative name server is unreachable, then the resolver may respond the client with old records. We wanted to measure how widespread this strategy is among resolvers. A DNS record, that is stored in a resolvers cache is considered stale, when that records TTL value is expired. Many resolvers clear the expired entries from their caches automatically to free the storage space. The deletion of stale records depends on the implementation of the cache data structure and the configuration of the resolver. The calculation of whether a record is stale in BIND9 only happens during the lookup dynamically. BIND9 does not know which lookups expire when and it uses a hash table to store and manage stale records.

Besides its advantages for static/long term IP addresses, stale DNS records can also cause issues with the resolution of domain names for IP addresses that change often, leading to errors or incorrect information being returned to users. If a DNS record is associated with an IP address that is no longer in use or has been reassigned to a different domain, the DNS system may still return the old IP address when a user attempts to access the domain. This can cause users to be directed to the wrong website or to receive an error message when trying to access the domain.

Recursive resolvers can be configured to delete cache entries that have been stored for a long time after a certain amount of time has passed. For example, BIND9 removes these entries after one week by default [1].

The goal of our first stale record experiment is to determine if a resolver is able to keep answering the client queries with stale records in case of total unavailability of the authoritative server, even though the resolvers cache entries are expired. This experiment also allows us to see if stale records can help in case the server has an outage.

Public recursive resolvers often use anycast to route traffic to multiple endpoint devices using a single destination address. This helps to distribute network traffic across multiple devices or locations, improving resilience and performance. With anycast, DNS queries are routed to the closest or most available DNS server, which can improve the speed and reliability of DNS resolution for users worldwide. However, the use of anycast can result in fragmented caches across multiple servers.

Each resolver IP address can have a different number of caches and the resolvers might be deployed behind a load balancer architecture to distribute the network load. The implementation of the load balancer influences the decision, which server should answer the incoming DNS query. With this architecture, each individual server might maintain a cache that is isolated from the other servers. To increase the cache hit rates, some architectures might use a shared cache pool that can be used by multiple servers[6].

3 Experiments and Results

Due to the varying number of caches among different recursive resolvers, we have to treat each resolver differently and be almost certain that we have first filled all the caches of a resolver with a resource record to be able to create stale record entries at the resolver side.

We divide our first stale record experiment in two phases, namely the prefetching phase and the stale phase. During the prefetching phase, we send multiple DNS queries to the DNS resolvers in order to increase the likelihood that all of the resolvers caches will be filled with a 95% chance, assuming the load balancers use a uniformly distributed implementation. We want to fill all the caches of the resolver with this resource record because if a cache is not hit, and we send the same DNS query to the resolver in the stale phase and hit an unfilled cache, the resolvers response will most likely be a resolution error. If we miss a resolvers cache, a stale record for that missed cache can't exist, therefore we cant test if the resolver supports stale records or not. One of the challenges of this phase is the unknown cache count of the resolvers. Every resolver can have different cache counts and these cache counts could be implemented in a way, so that it varies during the day based on the network load of the resolver. To estimate the number of caches for each DNS operator's resolvers, we conducted separate experiments at different times of the day and used the maximum cache count from these experiments as the result. Our results show that Cloudflare and OpenDNS have the most amount of caches, with 18 each. Some operators had varying numbers of caches throughout the day. Once we determined the number of caches for all resolvers, we began our stale record experiment using the identified cache amounts for each resolver.

To determine if a resolver is actually sending an outdated record rather than fetching a new record from our server, we altered the IP addresses of the DNS responses on our server after the prefetching phase, so that we can examine the received response on the client side using this IP address, if the answer of the resolver was stale. (Diagram for this?) After populating the caches of all resolver IP addresses during the prefetching phase, we simulated the selected packetloss rates on our authoritative name server. With the simulated packetloss rate, we started making DNS requests to the resolvers again. We then examined the DNS responses and extracted the results of this experiment.

In each stale record experiment, we use a randomly generated string with a length of 3 that is included in the query names sent during the experiment. This ensures that subsequent experiments do not interfere with each other.

3.8 Stale Record Support Experiment Results

The first stale record experiment separated the selected resolvers we were investigating into two groups: those that frequently send stale records and those that do not send stale records. We then compared the results between these two groups to see how the stale records affected DNS resolution.

During our experiments, we only received stale records from 6 of the 30 resolver IP addresses we examined. We noticed that a DNS operator might utilize stale records for

3 Experiments and Results

some of their IP addresses, while others may not use stale records at all. For example, we selected 3 resolver IP addresses belonging to the operator OpenDNS and found that 2 of them sent stale records, while the other did not. Our strategy of selecting one IP address per configuration proved useful, as different configurations may behave differently in the event of failure.

... (1. stale experiment plots)

When we compare the failure rates of these two groups, we clearly see that the query resolution process benefit from the stale records. While stale records can be useful for IP Addresses that doesn't change often, A stale DNS record is no longer accurate or relevant when a domain name is transferred to a new owner, when a website's IP address changes, or when a resource is no longer available.

We conclude that under normal operations (with authoritatives responding), recursive resolvers do not serve stale results (as has been proposed when the authoritative cannot be reached [7]).

3.9 Stale Record TTL Experiment

After the first experiment with stale records, we wanted to determine how long a resolver will send stale records, if it is able to send them at all. By answering this question, we can estimate how long a resolver can provide DNS responses as if there is no failure on the name server. In this second stale record experiment, we also want to examine the impact of the TTL value on the rate of stale records. Since it is not feasible to experiment with all possible TTL configurations, we chose a few TTL values to test, which are 60, 300, 900, and 3600. (Frequently used TTL values are 60, 1800, 3600, 86400) After the prefetching phase, we simulated 100% packet loss on our name server to determine if the resolvers would continue to answer client queries in the event of total unavailability of the name server. During the stale phase of the experiment, we sent queries every TTL second for a total of 6 hours.

3.10 Stale Record TTL Experiment Results

... (duration plots)

The results of the experiment indicate that the TTL values of the records may influence the duration of stale records. We found that resolvers in our experiment tended to send more stale records with lower TTL values. However, when the highest TTL value of 3600 was used, we never received a single stale record for 4 of the IP addresses, and OpenDNS rarely sent stale records.

(How resolvers alter the TTL, TTL stale records are either 30 or 0. A TTL value of 0 means that the DNS record is expired should not be cached and must be retrieved from the authoritative server every time it is needed. However, if a record should expire but

3 Experiments and Results

we receive a TTL value of 30, this could signal that the resolver refreshes this record by itself by adding additional time for this entry to be cached.)

It is possible that the resolvers check if a stale record is being requested in order to determine whether to retain the stale record in the cache because they are not able to retain that record indefinitely. It is likely that stale records that are frequently requested are kept in the resolver cache for a longer period of time.

3.11 Stale Record Send Frequency Experiment

To further narrow down how the stale record rates are affected, we conducted another experiment with a constant query sending frequency and varying TTL values. We used the same TTL values in the previous experiment and sent a query to each resolver every 2 minutes instead of waiting for the TTL interval after each query.

3.12 Stale Record Send Frequency Experiment Results

While the results suggest that frequency of queries may influence stale record rates, they do not clearly indicate whether this effect leads to higher or lower rates. While the OpenDNS resolvers had clearly higher stale record rates with lower query sending rates, the Dyn operator began to send SERVFAIL responses very early in the experiment when the query rate was lower.

3.13 Ripe Atlas Stale Record Experiment

In order to determine the prevalence of stale records among the RIPE Atlas probe resolvers, we conducted the same stale record experiment using 750 RIPE Atlas probes worldwide.

3.14 Ripe Atlas Stale Record Experiment Results

3.15 DNS Truncation Experiment

IP spoofing is a type of cyber attack in which an attacker alters the source IP address of a packet to make it appear as if it came from a different device. This is typically done to hide the identity of the attacker, or to launch attacks on networks or individual devices. "IP-Spoofing" in the context of networking means that the communication appears to originate from another host by manipulating IP packets with a forged sender IP address. DNS Cookies work by adding an additional layer of protection to the DNS protocol, allowing DNS servers and resolvers to authenticate each other and verify the integrity of DNS responses. DNS cookies use cryptographic techniques to create a small piece of data that is sent between a DNS resolver and a DNS server. Although they are not a complete solution for all security problems, they provide a lightweight mitigation

3 Experiments and Results

mechanism against attacks such as reflection attacks and DDoS attacks from spoofed sources. In order to benefit from DNS Cookies and decrease their vulnerability to the threats, both the client and server must support/implement it. Another way to protect against IP-Spoofing attacks is by using TCP.

TCP uses a 3-way handshake to initiate a connection between the sender and receiver, which uses a random number that prevents IP-Spoofing. Due to the sequence number mechanism of the TCP connections, this IP address manipulation via TCP on an existing connection is only possible by bruteforcing the sequence number or eavesdropping on the client, which makes TCP much more secure than UDP. UDP has no sequence numbers, therefore it is possible to insert manipulated data that be treated as if they were coming from the real host.

DNS queries are usually sent via the UDP protocol and as a result, it is possible to send a forged DNS query with a spoofed IP address to a resolver.

In the truncation experiment, we wanted to see what would happen if the server dictates that from now on it will only be addressed completely via TCP. This is important to ensure that the resolver always gets answers. For this purpose, we modified our iptables rule to only simulate packetloss on UDP packets and not TCP. We used a response policy zone in BIND9 to be able to answer a DNS query sent via UDP with a DNS response packet, where the truncation bit of the packet is set and there is no answer section in the response. TCP Packetloss may cause issues with 3-way handshake when the packetloss rate is very high, therefore we only simulated packetloss on UDP packets.

(If a DNS server is attacked or similar, it may be that these UDP packages come from spoofed IP addresses. This is not entirely visible to auth server.)

The truncation bit is a flag in the DNS header that indicates whether the response to a DNS query has been truncated due to size constraints. The truncation bit is used to allow DNS clients and servers to handle responses that are too large to fit in a single DNS message. If a DNS response is larger than the maximum size of a DNS message, the server can set the truncation bit in the response to indicate that the response has been truncated and that the client should send a new query using a different mechanism, such as TCP, to retrieve the full response.

A response policy zone (RPZ) in BIND9 is a feature that allows the DNS server to modify its responses to DNS queries based on a set of rules. The RPZ zone is a special type of zone that contains the rules for modifying DNS responses. These rules can be used to block access to certain domains, redirect traffic to different domains, or provide alternate responses for specific queries. The RPZ zone can be stored in a separate file on the DNS server, or it can be included in the server's main configuration file. Once the RPZ zone is configured, the DNS server will apply the rules in the RPZ zone to all DNS queries that it receives. If a query matches one of the rules in the RPZ zone, the DNS server will modify its response accordingly. Using an RPZ can potentially affect the latency of DNS queries, as the DNS server must perform additional processing when handling a query for a domain that has an associated RPZ policy. This additional processing may increase the time it takes for the DNS server to resolve the query and return a response

to the client.

(Paper notes: what when cookies enforced, support in world...)

3.16 DNS Truncation Experiment Results

All of the resolvers in the experiment were able to switch to TCP and send queries and responses via TCP. The results of the truncation experiments show close similarity to the packetloss experiment, with the addition of TCP packets. Our packetloss experiment showed that none of the resolver we selected switched their communication channel from UDP to TCP by themselves. Only with truncation experiment we saw TCP packets.

4 Conclusions

To summarize the main findings and implications of the research; If the caches of a resolver are full and the resolver supports stale records, these stale records help reduce the failure rate when the authoritative server is unreachable. The serving stale strategy is good for long term/static IP addresses while IP addresses that change often does not benefit from this. Stale record not widely supported for resolvers we tested. DNS retransmissions amplify traffic and increase resolution latency up to x times but can be helpful to resolve queries. Using TCP as an underlying protocol prevents IP spoofing, switching from UDP to TCP can be useful when a name server detects high volume of dns traffic to filter spoofed IP addresses. If there is too much packetloss on the server side, resolvers might detect this issue and give up the resolution process early, responding with a SERVFAIL.

While many websites use short cache timeouts to support control with DNS based load balancing, they may wish to consider longer timeouts as part of strategies for DDoS defense.

5 Discussion

Only PC users tested, mobile/other devices? Not all resolvers are tested, more TTL Values can be tested. Results show only the results of today, tomorrow might vary depending on implementation. Bei Ripe atlas keine Kontrolle darüber, was für Geräte wir haben, gewisse bias an ripe probes. eine Probe pro ASN gewaehlt.

Although we claim our study captures diverse recursive resolvers, we do not claim they are representative of a typical Internet client. ...

6 Ethical Considerations

All the experiments we introduced in this paper are performed in a benign manner so that our experiments does not cause loads of network traffic on the resolver side. The highest frequency that the queries are sent to a resolver was the prefetching phases of the stale records experiments. In the prefetching phase, we have sent 2 queries every second to a resolver with a high amount of cache count for about 30 seconds to be able to hit all the caches of the resolver in a short time. (In dns scan experiments, some parts of IP ranges are skipped due to privacy ...)

Bibliography

- [1] ISC BIND. 2020. BIND 9 Configuration Reference. <https://ftp.isc.org/isc/bind9/9.16.7/doc/arm/> (2020).
- [2] Cloudflare. 2022. The Internet’s Fastest, Privacy-First DNS Resolver. <https://1.1.1.1/> (2022).
- [3] John Heidemann Wouter B. de Vries Moritz Müller Lan Wei Giovane C. M. Moura, Ricardo de O. Schmidt and Christian Hesselman. 2016. Anycast vs. DDoS: Evaluating the November 2015 Root DNS Event. <https://doi.org/10.1145/2987443.2987446> (2016).
- [4] Moritz Müller Ricardo de O. Schmidt Marco Davids Giovane C. M. Moura, John Heidemann. 2018. When the Dike Breaks: Dissecting DNS Defenses During DDoS. (2018).
- [5] Google. [n. d.]. Public DNS. <https://developers.google.com/speed/public-dns/> ([n. d.]).
- [6] Google. 2021. Public DNS. <https://developers.google.com/speed/public-dns/docs/performance> (2021).
- [7] D. Lawrence and W. Kumari. 2017. Serving Stale Data to Improve DNS Resiliency-02. <https://www.ietf.org/archive/id/draft-tale-dnsop-serve-stale-02.txt> (2017).
- [8] OpenDNS. [n. d.]. Setup Guide: OpenDNS. <https://www.opendns.com/setupguide> ([n. d.]).
- [9] Quad9. [n. d.]. Internet Security & Privacy In a Few Easy Steps. <https://quad9.net> ([n. d.]).

Eidesstattliche Versicherung

(Affidavit)

Name, Vorname
(surname, first name)

Matrikelnummer
(student ID number)

☐ Bachelorarbeit
(Bachelor's thesis)

☐ Masterarbeit
(Master's thesis)

Titel
(Title)

Ich versichere hiermit an Eides statt, dass ich die vorliegende Abschlussarbeit mit dem oben genannten Titel selbstständig und ohne unzulässige fremde Hilfe erbracht habe. Ich habe keine anderen als die angegebenen Quellen und Hilfsmittel benutzt sowie wörtliche und sinngemäße Zitate kenntlich gemacht. Die Arbeit hat in gleicher oder ähnlicher Form noch keiner Prüfungsbehörde vorgelegen.

I declare in lieu of oath that I have completed the present thesis with the above-mentioned title independently and without any unauthorized assistance. I have not used any other sources or aids than the ones listed and have documented quotations and paraphrases as such. The thesis in its current or similar version has not been submitted to an auditing institution before.

Ort, Datum
(place, date)

Unterschrift
(signature)

Belehrung:

Wer vorsätzlich gegen eine die Täuschung über Prüfungsleistungen betreffende Regelung einer Hochschulprüfungsordnung verstößt, handelt ordnungswidrig. Die Ordnungswidrigkeit kann mit einer Geldbuße von bis zu 50.000,00 € geahndet werden. Zuständige Verwaltungsbehörde für die Verfolgung und Ahndung von Ordnungswidrigkeiten ist der Kanzler/die Kanzlerin der Technischen Universität Dortmund. Im Falle eines mehrfachen oder sonstigen schwerwiegenden Täuschungsversuches kann der Prüfling zudem exmatrikuliert werden. (§ 63 Abs. 5 Hochschulgesetz - HG -).

Die Abgabe einer falschen Versicherung an Eides statt wird mit Freiheitsstrafe bis zu 3 Jahren oder mit Geldstrafe bestraft.

Die Technische Universität Dortmund wird ggf. elektronische Vergleichswerkzeuge (wie z.B. die Software „turnitin“) zur Überprüfung von Ordnungswidrigkeiten in Prüfungsverfahren nutzen.

Die oben stehende Belehrung habe ich zur Kenntnis genommen:

Official notification:

Any person who intentionally breaches any regulation of university examination regulations relating to deception in examination performance is acting improperly. This offense can be punished with a fine of up to EUR 50,000.00. The competent administrative authority for the pursuit and prosecution of offenses of this type is the Chancellor of TU Dortmund University. In the case of multiple or other serious attempts at deception, the examinee can also be unenrolled, Section 63 (5) North Rhine-Westphalia Higher Education Act (*Hochschulgesetz, HG*).

The submission of a false affidavit will be punished with a prison sentence of up to three years or a fine.

As may be necessary, TU Dortmund University will make use of electronic plagiarism-prevention tools (e.g. the "turnitin" service) in order to monitor violations during the examination procedures.

I have taken note of the above official notification:*

Ort, Datum
(place, date)

Unterschrift
(signature)

***Please be aware that solely the German version of the affidavit ("Eidesstattliche Versicherung") for the Bachelor's/ Master's thesis is the official and legally binding version.**